

```
In [1]: #!/pip install tensorflow_hub
```

```
In [2]: import tensorflow_hub as hub
import tensorflow as tf
from matplotlib import pyplot as plt
import numpy as np
import cv2
import os #navigate path structure
from PIL import Image
```

```
In [3]: model = hub.load('https://tfhub.dev/google/magenta/arbitrary-image-stylization-v1-2')
```

```
In [4]: def load_image(img_path):
    img = tf.io.read_file(img_path)
    img = tf.image.decode_image(img, channels=3)
    img = tf.image.convert_image_dtype(img, tf.float32)
    img = img[tf.newaxis, :]
    return img
```

## Sample 1

```
In [5]: content_image = load_image('C:/Users/soumi/3.DMML_2/dataset/content/251 (3).jpg')
style_image = load_image('C:/Users/soumi/3.DMML_2/dataset/art_work/Henri_Matisse_9')
```

```
In [6]: content_image.shape
```

```
Out[6]: TensorShape([1, 533, 800, 3])
```


```
In [7]: style_image.shape
```

```
Out[7]: TensorShape([1, 723, 542, 3])
```

```
In [8]: plt.imshow(np.squeeze(content_image))
plt.show()
```




```
In [9]: plt.imshow(np.squeeze(style_image))  
plt.show()
```



```
In [10]: stylized_image = model(tf.constant(content_image), tf.constant(style_image))[0]  
stylized_image1 = stylized_image
```

```
In [11]: plt.imshow(np.squeeze(stylized_image))  
plt.show()
```



```
In [12]: cv2.imwrite('generated_img.jpg', cv2.cvtColor(np.squeeze(stylized_image)*255, cv2.COLOR_RGB2BGR))
```

Out[12]: True


```
In [13]: # Display the images using matplotlib
plt.figure(figsize=(12, 8))

plt.subplot(1, 3, 1)
plt.title('Content Image')
plt.imshow(np.squeeze(content_image))
plt.axis('off')

plt.subplot(1, 3, 2)
plt.title('Style Image')
plt.imshow(np.squeeze(style_image))
plt.axis('off')

plt.subplot(1, 3, 3)
plt.title('Stylized Image')
plt.imshow(np.squeeze(stylized_image))
plt.axis('off')

plt.show()
```



```
In [14]: # Resize the style_image to match the shape of the content_image
style_image_resized = tf.image.resize(style_image, content_image.shape[1:3], method=

# Calculate Mean Squared Error
mse1 = tf.reduce_mean(tf.square(content_image - style_image_resized)).numpy()
print(f"Mean Squared Error: {mse1}")
```

Mean Squared Error: 0.1771637201309204

```
In [15]: vgg = tf.keras.applications.VGG16(include_top=False, weights="imagenet")

def preprocess_for_vgg(image):
    return tf.keras.applications.vgg16.preprocess_input(image * 255)

content_image_vgg = preprocess_for_vgg(content_image)
stylized_image_vgg = preprocess_for_vgg(stylized_image)

content_features = vgg(content_image_vgg)
stylized_features = vgg(stylized_image_vgg)

perceptual_loss1 = tf.reduce_mean(tf.square(content_features - stylized_features))
print(f"Perceptual Loss: {perceptual_loss1}")
```

Perceptual Loss: 73.82260131835938

```
In [16]: def gram_matrix(feature_maps):
    num_channels = feature_maps.shape[-1]
    flattened = tf.reshape(feature_maps, (-1, num_channels))
    gram = tf.matmul(flattened, flattened, transpose_a=True)
    gram /= tf.cast(tf.reduce_prod(feature_maps.shape), tf.float32)
    return gram

# Assuming you have already defined vgg and preprocess_for_vgg

style_image_vgg = preprocess_for_vgg(style_image)
stylized_image_vgg = preprocess_for_vgg(stylized_image) # You might have missed to

style_features = vgg(style_image_vgg)
stylized_features = vgg(stylized_image_vgg)

style_gram = gram_matrix(style_features)
stylized_gram = gram_matrix(stylized_features)

gram_loss1 = tf.reduce_mean(tf.square(style_gram - stylized_gram)).numpy()
print(f"Gram Matrix Loss: {gram_loss1}")
```

Gram Matrix Loss: 0.0008529422921128571



```
In [17]: def total_variation_loss(image):
          x_deltas = image[:, :, 1:, :] - image[:, :, :-1, :]
          y_deltas = image[:, 1:, :, :] - image[:, :-1, :, :]
          tv_loss = tf.reduce_sum(tf.abs(x_deltas)) + tf.reduce_sum(tf.abs(y_deltas))
          return tv_loss

          # Calculate total variation loss for the stylized image
          tv_loss1 = total_variation_loss(stylized_image)
          print(f"Total Variation Loss: {tv_loss1.numpy()}")
```

Total Variation Loss: 67402.3125

## Sample 2

```
In [18]: content_image = load_image('C:/Users/soumi/3.DMML_2/dataset/content/391.jpg')
          style_image = load_image('C:/Users/soumi/3.DMML_2/dataset/art_work/Claude_Monet_56')
```

```
In [19]: content_image.shape
```

```
Out[19]: TensorShape([1, 1367, 2050, 3])
```


```
In [20]: style_image.shape
```

```
Out[20]: TensorShape([1, 1379, 2560, 3])
```

```
In [21]: plt.imshow(np.squeeze(content_image))
          plt.show()
```




```
In [22]: plt.imshow(np.squeeze(style_image))
          plt.show()
```



```
In [23]: stylized_image = model(tf.constant(content_image), tf.constant(style_image))[0]
stylized_image2 = stylized_image
```

```
In [24]: plt.imshow(np.squeeze(stylized_image))
plt.show()
```



```
In [25]: cv2.imwrite('generated_img.jpg', cv2.cvtColor(np.squeeze(stylized_image)*255, cv2.COLOR_RGB2BGR))
```

```
Out[25]: True
```

```
In [26]: # Display the images using matplotlib
plt.figure(figsize=(12, 8))

plt.subplot(1, 3, 1)
plt.title('Content Image')
plt.imshow(np.squeeze(content_image))
plt.axis('off')
```

```
plt.subplot(1, 3, 2)
plt.title('Style Image')
plt.imshow(np.squeeze(style_image))
plt.axis('off')

plt.subplot(1, 3, 3)
plt.title('Stylized Image')
plt.imshow(np.squeeze(stylized_image))
plt.axis('off')

plt.show()
```



```
In [27]: # Resize the style_image to match the shape of the content_image
style_image_resized = tf.image.resize(style_image, content_image.shape[1:3], method=

# Calculate Mean Squared Error
mse2 = tf.reduce_mean(tf.square(content_image - style_image_resized)).numpy()
print(f"Mean Squared Error: {mse2}")
```

Mean Squared Error: 0.2287299931049347

```
In [28]: content_image_vgg = preprocess_for_vgg(content_image)
stylized_image_vgg = preprocess_for_vgg(stylized_image)

content_features = vgg(content_image_vgg)
stylized_features = vgg(stylized_image_vgg)

perceptual_loss2 = tf.reduce_mean(tf.square(content_features - stylized_features))
print(f"Perceptual Loss: {perceptual_loss2}")
```

Perceptual Loss: 44.7812385559082

```
In [29]: style_image_vgg = preprocess_for_vgg(style_image)
stylized_image_vgg = preprocess_for_vgg(stylized_image) # You might have missed to

style_features = vgg(style_image_vgg)
stylized_features = vgg(stylized_image_vgg)

style_gram = gram_matrix(style_features)
stylized_gram = gram_matrix(stylized_features)

gram_loss2 = tf.reduce_mean(tf.square(style_gram - stylized_gram)).numpy()
print(f"Gram Matrix Loss: {gram_loss2}")
```

Gram Matrix Loss: 3.1323754228651524e-05

```
In [30]: # Calculate total variation loss for the stylized image
tv_loss2 = total_variation_loss(stylized_image)
print(f"Total Variation Loss: {tv_loss2.numpy()}")
```

Total Variation Loss: 406243.75

In [ ]:

## Sample 3

```
In [52]: content_image = load_image('C:/Users/soumi/3.DMML_2/dataset/content/56.jpg')  
style_image = load_image('C:/Users/soumi/3.DMML_2/dataset/art_work/Eugene_Delacroix')
```

```
In [53]: content_image.shape
```

```
Out[53]: TensorShape([1, 1600, 2560, 3])
```

```
In [54]: style_image.shape
```

```
Out[54]: TensorShape([1, 600, 735, 3])
```

```
In [55]: plt.imshow(np.squeeze(content_image))  
plt.show()
```



```
In [56]: plt.imshow(np.squeeze(style_image))  
plt.show()
```





```
In [57]: stylized_image = model(tf.constant(content_image), tf.constant(style_image))[0]
stylized_image3 = stylized_image
```

```
In [58]: plt.imshow(np.squeeze(stylized_image))
plt.show()
```



```
In [59]: cv2.imwrite('generated_img.jpg', cv2.cvtColor(np.squeeze(stylized_image)*255, cv2.COLOR_RGB2BGR))
```

```
Out[59]: True
```

```
In [60]: # Display the images using matplotlib
plt.figure(figsize=(12, 8))
```

```
plt.subplot(1, 3, 1)
plt.title('Content Image')
plt.imshow(np.squeeze(content_image))
plt.axis('off')

plt.subplot(1, 3, 2)
plt.title('Style Image')
plt.imshow(np.squeeze(style_image))
plt.axis('off')

plt.subplot(1, 3, 3)
plt.title('Stylized Image')
plt.imshow(np.squeeze(stylized_image))
plt.axis('off')

plt.show()
```



```
In [61]: # Resize the style_image to match the shape of the content_image
style_image_resized = tf.image.resize(style_image, content_image.shape[1:3], method=

# Calculate Mean Squared Error
mse3 = tf.reduce_mean(tf.square(content_image - style_image_resized)).numpy()
print(f"Mean Squared Error: {mse3}")
```

Mean Squared Error: 0.13691847026348114

```
In [62]: content_image_vgg = preprocess_for_vgg(content_image)
stylized_image_vgg = preprocess_for_vgg(stylized_image)

content_features = vgg(content_image_vgg)
stylized_features = vgg(stylized_image_vgg)

perceptual_loss3 = tf.reduce_mean(tf.square(content_features - stylized_features))
print(f"Perceptual Loss: {perceptual_loss3}")
```

Perceptual Loss: 44.0880126953125

```
In [63]: style_image_vgg = preprocess_for_vgg(style_image)
stylized_image_vgg = preprocess_for_vgg(stylized_image) # You might have missed to

style_features = vgg(style_image_vgg)
stylized_features = vgg(stylized_image_vgg)

style_gram = gram_matrix(style_features)
stylized_gram = gram_matrix(stylized_features)

gram_loss3 = tf.reduce_mean(tf.square(style_gram - stylized_gram)).numpy()
print(f"Gram Matrix Loss: {gram_loss3}")
```

Gram Matrix Loss: 0.0013236101949587464

```
In [64]: # Calculate total variation loss for the stylized image
tv_loss3 = total_variation_loss(stylized_image)
```

```
print(f"Total Variation Loss: {tv_loss3.numpy()}")
```

Total Variation Loss: 852280.375

In [ ]:

## Sample 4

```
In [66]: content_image = load_image('C:/Users/soumi/3.DMML_2/dataset/content/35 (2).jpg')  
style_image = load_image('C:/Users/soumi/3.DMML_2/dataset/art_work/Andy_Warhol_24..')
```

```
In [67]: content_image.shape
```

```
Out[67]: TensorShape([1, 1000, 1500, 3])
```

```
In [68]: style_image.shape
```


```
Out[68]: TensorShape([1, 676, 672, 3])
```

```
In [69]: plt.imshow(np.squeeze(content_image))  
plt.show()
```




```
In [70]: plt.imshow(np.squeeze(style_image))  
plt.show()
```





```
In [71]: stylized_image = model(tf.constant(content_image), tf.constant(style_image))[0]
stylized_image4 = stylized_image
```

```
In [72]: plt.imshow(np.squeeze(stylized_image))
plt.show()
```



```
In [73]: cv2.imwrite('generated_img.jpg', cv2.cvtColor(np.squeeze(stylized_image)*255, cv2.COLOR_RGB2BGR))
```

```
Out[73]: True
```

```
In [74]: # Display the images using matplotlib
plt.figure(figsize=(12, 8))
```



```
plt.subplot(1, 3, 1)
plt.title('Content Image')
plt.imshow(np.squeeze(content_image))
plt.axis('off')

plt.subplot(1, 3, 2)
plt.title('Style Image')
plt.imshow(np.squeeze(style_image))
plt.axis('off')

plt.subplot(1, 3, 3)
plt.title('Stylized Image')
plt.imshow(np.squeeze(stylized_image))
plt.axis('off')

plt.show()
```



```
In [75]: # Resize the style_image to match the shape of the content_image
style_image_resized = tf.image.resize(style_image, content_image.shape[1:3], method='bilinear')

# Calculate Mean Squared Error
mse4 = tf.reduce_mean(tf.square(content_image - style_image_resized)).numpy()
print(f"Mean Squared Error: {mse4}")
```

Mean Squared Error: 0.25097450613975525

```
In [76]: content_image_vgg = preprocess_for_vgg(content_image)
stylized_image_vgg = preprocess_for_vgg(stylized_image)

content_features = vgg(content_image_vgg)
stylized_features = vgg(stylized_image_vgg)

perceptual_loss4 = tf.reduce_mean(tf.square(content_features - stylized_features))
print(f"Perceptual Loss: {perceptual_loss4}")
```

Perceptual Loss: 171.49356079101562

```
In [77]: style_image_vgg = preprocess_for_vgg(style_image)
stylized_image_vgg = preprocess_for_vgg(stylized_image) # You might have missed to

style_features = vgg(style_image_vgg)
stylized_features = vgg(stylized_image_vgg)

style_gram = gram_matrix(style_features)
stylized_gram = gram_matrix(stylized_features)

gram_loss4 = tf.reduce_mean(tf.square(style_gram - stylized_gram)).numpy()
print(f"Gram Matrix Loss: {gram_loss4}")
```

Gram Matrix Loss: 0.000519386085215956

```
In [78]: # Calculate total variation loss for the stylized image
tv_loss4 = total_variation_loss(stylized_image)
print(f"Total Variation Loss: {tv_loss4.numpy()}")
```

Total Variation Loss: 277654.90625

## Comparing MSE


```
In [104... import matplotlib.pyplot as plt

# List of image names and their corresponding MSE values
image_names = ['stylized_image1', 'stylized_image2', 'stylized_image3', 'stylized_
mse_values = [mse1, mse2, mse3, mse4, ] # Replace with actual MSE values

# Create a bar chart
plt.bar(image_names, mse_values, color=['blue', 'orange', 'green'])

# Add labels and title
plt.xlabel('Image')
plt.ylabel('Mean Squared Error (MSE)')
plt.title('Comparison of Mean Squared Error (MSE) for Different Images')

# Display the chart
plt.show()
```



In [ ]: